

speedshop



Determining Ruby Process Counts

speedshop

THE COMPLETE GUIDE TO RAILS PERFORMANCE

NATE BERKOPEC

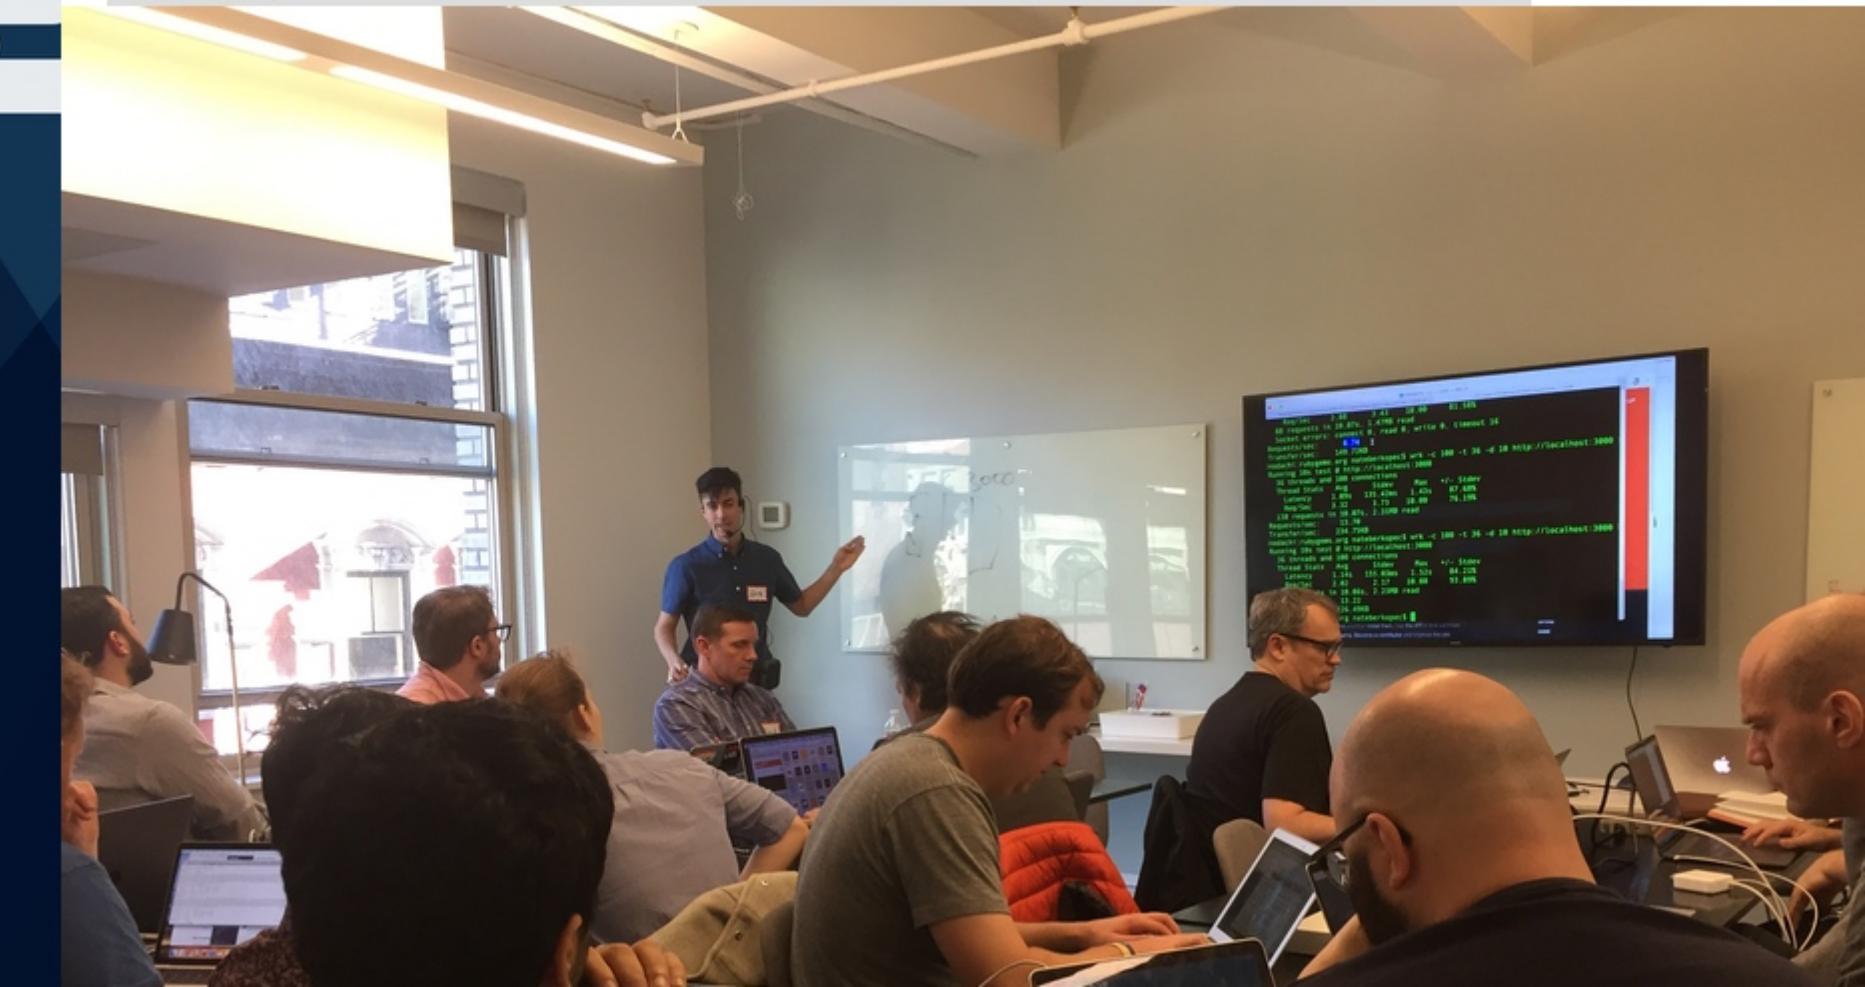
3 ActiveRecord Mistakes That Slow Down Rails Apps: Count, Where and Present

by Nate Berkoperc (@nateberkopec) of **speedshop** ([who?](#)), a Rails performance consultancy.

Summary: Many Rails developers don't understand what causes ActiveRecord to actually execute a SQL query. Let's look at three common cases: misuse of the count method, using where to select subsets, and the present? predicate. You may be causing extra queries and N+1s through the abuse of these three methods. (2778 words / 12 minutes)

SHARE: [f](#)

E-Mail





OUTLINE

1. Understanding the math
2. What the math means
3. Math Check!
4. Real world application

**HOW MANY
SERVERS DO YOU NEED?**

**300 REQUESTS PER
MINUTE**

250 MILLISECOND

**AVERAGE
1 SECOND 95TH
PERCENTILE**

LITTLE'S LAW

ITEMS IN SYSTEM =
ARRIVAL RATE X TIME IN
SYSTEM



**WAIT TIMES
UTILIZATION
10 CUSTOMERS ENTER LINE PER
MINUTE
TAKES 1 MINUTE TO CHECKOUT
CUSTOMER
20 CHECKOUT COUNTERS**

AN EXAMPLE WITH SIMPLE NUMBERS

100 REQUESTS PER MINUTE

250 MILLISECOND RESPONSE

TIMES

5 SERVERS

5 WORKER PROCESSES PER

WHAT DOES IT MEAN

**TRUE
IN THE LONG RUN**

**4 95TH PERCENTILE REQUESTS IN A ROW
HAPPENS 1 IN 100,000 REQUESTS
3 TIMES PER DAY AT 200 RPM**

UNDERSTANDING REQUEST QUEUES IN RUBY WEB APPS

**SCALING REDUCES WAIT TIMES
NOT RESPONSE TIMES**

**AUTOSCALING
MUST BE BASED ON WAIT TIMES**

HOW TO FIND REQUEST QUEUE TIMES

EASY TO SEE UNDERSCALING
HARD TO SEE OVERSCALING

**PREDICTABLE RESPONSE TIMES
ARE SCALABLE RESPONSE TIMES**

STORY

REDUCING WAIT TIMES ON A HIGH 95-TH APP

WHAT WE KNOW SO FAR

WIP = LATENCY * ARRIVAL RATE

SCALING REDUCES QUEUE

PREDICTABLE IS SCALABLE

HOW UNICORN WORKS

TWITTER, 2008

- » 600 requests/second
- » 180 application processes (mongrel)
- » About 300ms average server response time

SHOPIFY, 2013

- » Shopify receives 833 requests/second.
- » They average a 72ms response time
- » They run 53 application servers with a total of 1172 application processes (!!!) with NGINX and Unicorn.

ENVATO, 2013

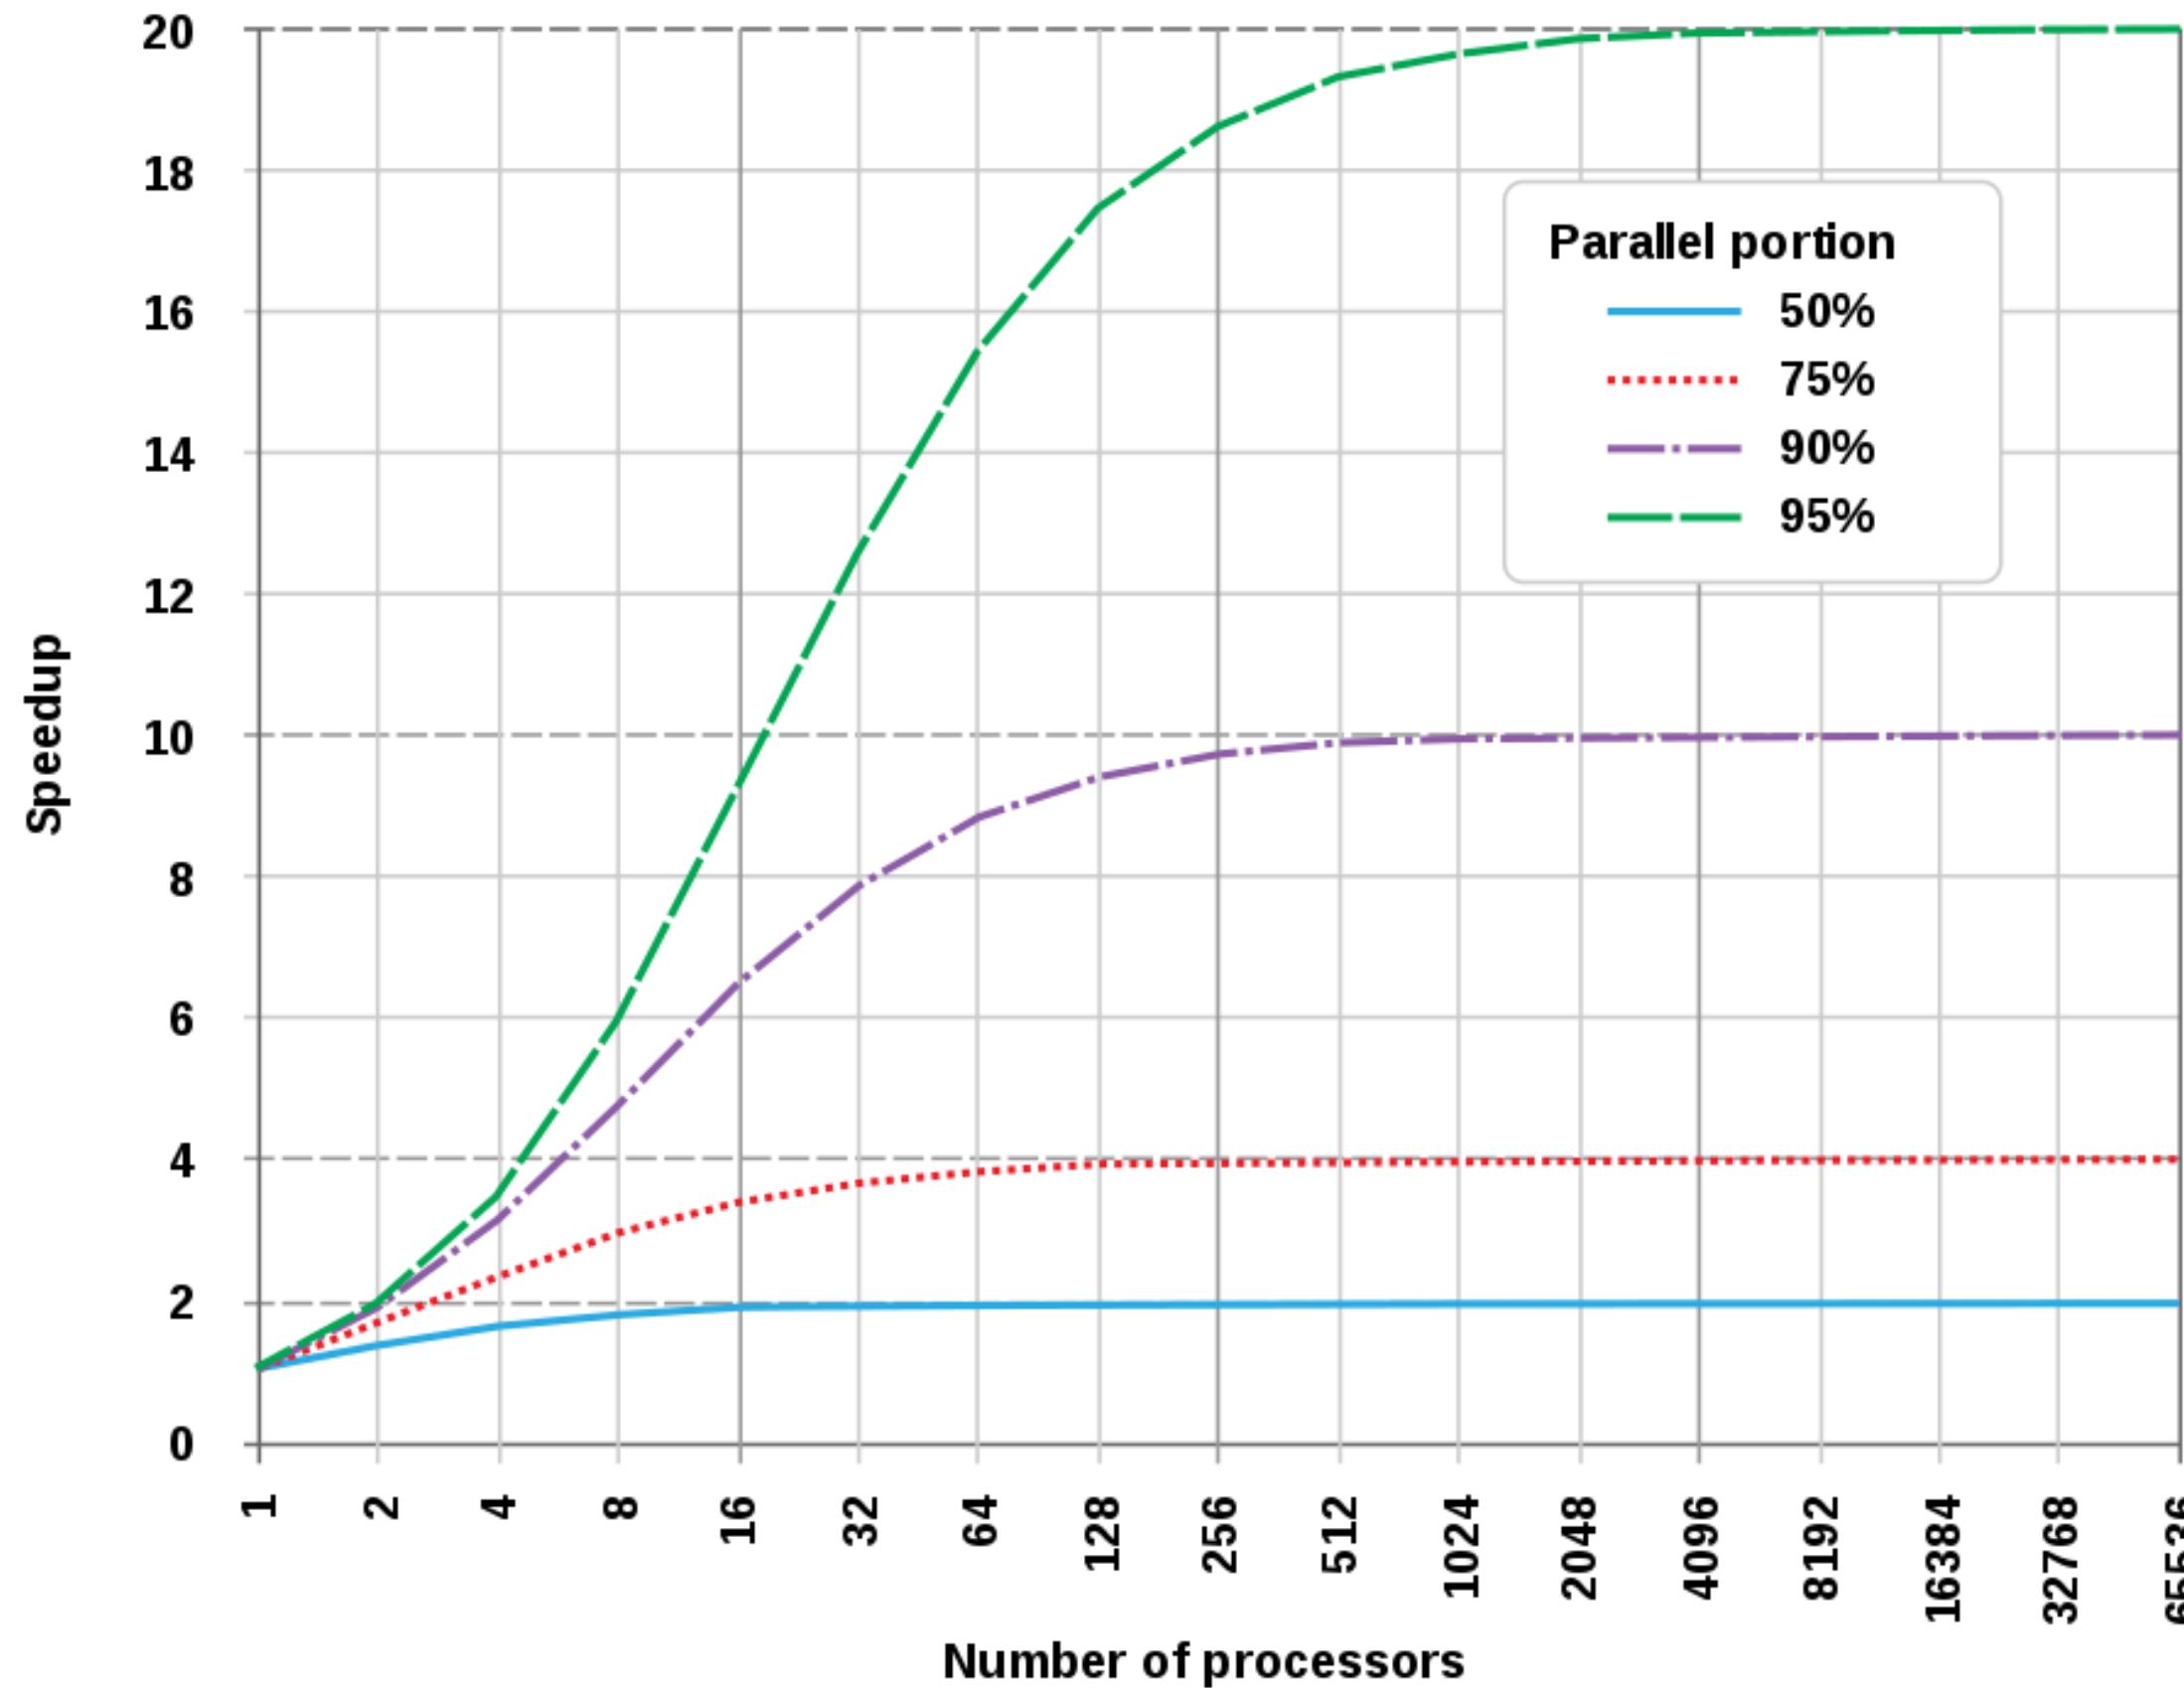
- » Envato receives 115 requests per second
- » They run an average of 147ms response time
- » They run 45 processes

WHY THE PROCESS?

THREADS VS PROCESSES

AMDAHL'S LAW

Amdahl's Law



g

h

w

z

WHAT % SPENT HOLDING GVL?

BREAKING THE LAWS

DIVIDING SERVERS UP - BIG OR SMALL?

WHERE DO REQUESTS QUEUE?

LOAD BALANCER BEHAVIOR

PASSENGER BEHAVIOR (UNIQUE CASE)

FILLING UP RESOURCES (CPU/MEM)

SERVERS

CONTAINER STORY

**OTHER SERVICES
BACKGROUND JOBS**

**OTHER SERVICES
DB, JRUBY, LANGUAGES**

CHECKIN: WHAT WE KNOW

LITTLE'S LAW

HOW REQUEST QUEUES WORK

HOW TO PROVISION

HOW TO SAVE \$\$\$

THANK YOU!

@NATEBERKOPEC

MORE TACTICS AT:

RAILSSPEED.COM